



# Additive models

Manel Martínez-Ramón

ECE, UNM

December 4, 2019

A Classification and Regression Tree (CART) is a recursive binary representation of the data




Root nodes contain a single feature, either  $X_1$  or  $X_2$ . Leaf nodes contain partitions of an  $\mathbb{R}^2$  space.

If  $\mathbf{v}_m : \{X_m \leq t_m\}$ , then the response of the tree can be coded as a linear combination of basis functions  $\phi(\cdot)$


$$f(\mathbf{x}) = \mathbb{E}[y|\mathbf{x}] = \sum_{m=1}^M w_m \mathbb{I}(\mathbf{x} \in R_m) = \sum_{m=1}^M w_m \phi(\mathbf{x}; \mathbf{v}_m)$$

where  $R_m$  is a region and  $w_m$  is the mean response in this region.

## Classification trees.



	D features (attributes)			Label
	Color	Shape	Size (cm)	
	Blue	Square	10	1
	Red	Ellipse	2.4	1
	Red	Ellipse	20.7	0



Root nodes contain a single feature. Leaf nodes contain values  $[n_0, n_1]$  of instances of each class matching the feature value.

- ❶ function fitTree(node,  $\mathcal{D}$ , depth);
- ❷ node.prediction = mean( $y_i : i \in \mathcal{D}$ ) // or class label distribution;
- ❸  $(j^*, t^*, \mathcal{D}_L, \mathcal{D}_R) = \text{split}(\mathcal{D})$ ;
- ❹ if not *worthSplitting*(depth, cost,  $\mathcal{D}_L$ ,  $\mathcal{D}_R$ ) then
- ❺     return node
- ❻ else
- ❼     node.test =  $\lambda \mathbf{x}. x_{j^*} < t^*$  // anonymous function;
- ❽     node.left = fitTree(node,  $\mathcal{D}_L$ , depth+1);
- ❾     node.right = fitTree(node,  $\mathcal{D}_R$ , depth+1);
- ❿     return node;

$\mathcal{D}$  refers to the training data set.

The split function chooses the best feature  $x_{ij}$  and the best value  $t$  according to a given cost function:

$$(j^*, t^*) = \arg \min_{j \in \{1 \dots D\}} \min_{t \in \mathcal{T}_j} \{ \text{cost}(\{\mathbf{x}_i, y_i; x_{ij} \leq t\}) + \text{cost}(\{\mathbf{x}_i, y_i; x_{ij} > t\}) \}$$

The *worthSplitting* function checks if a node is worth splitting using some heuristic:

- is the reduction in cost too small?
- has the tree exceeded the maximum desired depth?
- is the distribution of the response in either  $\mathcal{D}_L$  or  $\mathcal{D}_R$  sufficiently homogeneous (e.g., all labels are the same, so the distribution is pure)?
- is the number of examples in either  $\mathcal{D}_L$  or  $\mathcal{D}_R$  too small?

We define the cost as follows:

$$\text{cost}(\mathcal{D}) = \sum_{i \in \mathcal{D}} (y_i - \bar{y})^2$$

where  $\bar{y}$  is the mean of the response variable in the specified set of data. Alternatively, we can fit a linear regression model for each leaf, using as inputs the features that were chosen on the path from the root, and then measure the residual error.



First, a Multinoulli distribution of the current split must be estimated for all classes:

$$\hat{\pi}_c = p(y_i = c)$$

Then a probabilistic measure has to be computed:

- Probability of error:  $1 - \arg \max_c \hat{\pi}_c$
- Entropy:  $\mathbb{H}(\hat{\boldsymbol{\pi}}) = - \sum_{c=1}^C \hat{\pi}_c \log \hat{\pi}_c$
- Gini Index:  $\sum_{c=1}^C \hat{\pi}_c (1 - \hat{\pi}_c)$





Example of classification among three different classes of flowers (Iris database).

## Advantages:

- Easy to interpret;
- Handle mixed discrete and continuous inputs;
- Insensitive to monotone transformations of the inputs;
- Automatic variable selection,
- Relatively robust to outliers,
- Scale well to large data sets, and
- they can be modified to handle missing inputs

## Disadvantages:

- Not very accurate compared to other kinds of model.
- Trees are unstable: small changes to the input data can have large effects on the structure of the tree.

In order to reduce the variance of trees, the bagging technique averages many different estimates:

$$f(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M f_m(\mathbf{x})$$

where samples to train different trees are chosen randomly with repetition (Breiman, 1996).

Random forests use trees with different randomly chosen features and different random subsets (Breiman 2001)

A generalized additive model has the form

$$f(\mathbf{x}) = \alpha + \sum_{i=1}^D f_i(x_i)$$

where  $f(x_i) = \beta_i \phi(x_i)$ . The model can be applied to a cost function with a regularizer term

$$\sum_{i=1}^N e_i^2 + \sum_{j=1}^D \lambda_j \int \beta_j^2 \phi^2(t) dt$$

Interactions between variables can be added:

$$f(\mathbf{x}) = \alpha + \sum_{i=1}^D f_i(x_i) + \sum_{i,j} f_i(x_i, x_j) + \sum_{i,j,k} f_i(x_i, x_j, x_k) + \cdots$$

A particular case of this is the Volterra expansion of Module 5. In general, any type of spline functions can be used in this model.

Consider the additive model

$$f(\mathbf{x}) = \omega_0 + \sum_{m=1}^m \omega_m \phi_m(\mathbf{x})$$

and the exponential cost function

$$\exp(-y_i f(\mathbf{x}_i)), \quad y_i \in \{-1, 1\}$$

The minimization problem at step  $m$  is

$$L_m(\phi) = \sum_{i=1}^N \exp[-y_i (f_{m-1}(\mathbf{x}_i) + \beta_m \phi_m(\mathbf{x}_i))]$$

This can be rewritten as

$$\begin{aligned} L_m(\phi) &= \sum_{i=1}^N \exp [-y_i (f_{m-1}(\mathbf{x}_i) + \beta_m \phi_m(\mathbf{x}_i))] \\ &= \sum_{i=1}^N \omega_{i,m} \exp [-y_i \beta_m \phi_m(\mathbf{x}_i)] \end{aligned}$$

where

$$\begin{aligned} \omega_{i,m} &= \exp (-y_i (f_{m-1}(\mathbf{x}_i))) \\ &= \omega_{i,m-1} \exp (-y_i \beta_m (\phi_m(\mathbf{x}_i))) \end{aligned}$$

is a weight applied to sample  $i$  at iteration  $m$ .

The problem consists of minimizing

$$L_m(\phi) = \sum_{i=1}^N \omega_{i,m} \exp[-y_i \beta_m \phi_m(\mathbf{x}_i)]$$

Considering the convexity of the exponential:

$$\begin{aligned} & \sum_{i=1}^N \omega_{i,m} \exp[-y_i \beta_m \phi_m(\mathbf{x}_i)] \\ & \leq \sum_{i=1}^N \frac{1 - y_i \phi_m(\mathbf{x}_i)}{2} \omega_{i,m} \exp(\beta_m) + \sum_{i=1}^N \frac{1 + y_i \phi_m(\mathbf{x}_i)}{2} \omega_{i,m} \exp(-\beta_m) \\ & = e_m \exp(\beta_m) + (1 - e_m) \exp(-\beta_m) \end{aligned}$$

Nulling the derivative with respect to  $\beta_m$  gives its optimal value

$$\beta_m = \frac{1}{2} \log \frac{1 - e_m}{e_m}$$

where  $e_m = \sum_{i=1}^N \omega_{i,m} \mathbb{I}(y_i \neq \phi_m(\mathbf{x}_i))$



- ➊ Training data  $\{\mathbf{x}_1; y_1 \cdots \mathbf{x}_N; y_N\}$
- ➋ Set  $\omega_{i,0} = \frac{1}{M}$
- ➌ **for**  $m$  in  $1 \cdots M$ 
  - ➊ Train weak learner  $\phi_m$  with distribution  $\omega_{i,m}$
  - ➋ Compute the weighted error  $e_m = \sum_{i=1}^N \omega_{i,m} \mathbb{I}(y_i \neq \phi_m(\mathbf{x}_i))$
  - ➌ Compute  $\beta_m = \frac{1}{2} \log \frac{1-e_m}{e_m}$
  - ➍ Update weights  $\omega_{i,m} = \omega_{i,m-1} \exp(-y_i \beta_m (\phi_m(\mathbf{x}_i)))$
  - ➎ Normalize  $\omega_{i,m}$  so that  $\sum_i \omega_{i,m} = 1$
- ➍ Final hypothesis:

$$f(\mathbf{x}) = \text{sign} \left( \sum_{m=1}^M \beta_m \phi_m(\mathbf{x}) \right)$$

A simple way to train a learner  $\phi_m$  is to adjust a linear classifier using the distribution  $\omega_{i,m}$ . Thus, the learner is

$$\phi_m(\mathbf{x}_i) = \text{sign}(\mathbf{w}_m^\top \mathbf{x}_i)$$

where it is assumed that a bias is included in the weight vector. The optimization can be written as

$$\begin{aligned} \min_{\mathbf{w}} \sum_{i=1}^N \omega_{i,m} (y_i - \mathbf{w}_m^\top \mathbf{x}_i)^2 &= \min_w \|\mathbf{\Omega}_m^{\frac{1}{2}} \mathbf{y} - \mathbf{\Omega}_m^{\frac{1}{2}} \mathbf{X}^\top \mathbf{w}_m\|^2 \\ &= \min_{\mathbf{w}} \left( \mathbf{w}_m^\top \mathbf{X} \mathbf{\Omega}_m \mathbf{X} \mathbf{w}_m - \mathbf{w}_m^\top \mathbf{X} \mathbf{\Omega}_m \mathbf{y} \right) \end{aligned}$$

where  $\mathbf{\Omega}_m = \text{diag}[\omega_{1,m} \cdots \omega_{N,m}]$  and whose solution is

$$\mathbf{w}_m = \left( \mathbf{X} \mathbf{\Omega}_m \mathbf{X}^\top \right)^{-1} \mathbf{X} \mathbf{\Omega}_m \mathbf{y}$$

- Nevertheless, a better solution is to simply choose base learners  $\phi_m$  with VC dimension  $h = 2$ , this is, expressed in a subspace of one dimension.
- They can be seen as decision trees of one single node.

In order to train a tree, we must look for the optimal feature in  $\mathbf{x}$  and its optimal threshold. A way to do it is

- 1 **for**  $j$  in  $1 \cdots D$ 
  - 1 sort all instances of feature  $x_{i,j}$
  - 2 use them as threshold to split data  $\mathbf{x}_i$  and compute the split error
  - 3 Choose the threshold that produces the minimum error.
- 2 Choose the dimension  $j$  and threshold  $t_j$  that produce the minimum error.

- This lesson contains a summary of ensemble learning.
- The simplest adaptive basis function structure is the classification and regression tree (CART). In spite of its advantages, its performance is poor compared to other structures.
- Ensembles of random trees constructed with subsets of features are called random forests. They are known to have excellent performance at expenses of a considerable computational burden.
- An alternative ensemble learning methodology is boosting.
- The Adaboost algorithm combines weak learners. The learners are trained with the most difficult samples, and then they are weighted as a function of its quality.